

A Title: with a longer subtitle if needed

Author's Name Here

March 21, 2015

Abstract

This is where your abstract goes. An abstract should be a short paragraph telling the reader what you did, who might care, and why. Your abstract is what will be read first; most readers will use the abstract to decide whether or not to read the rest of the paper. It is therefore important to have an abstract that is both clear and concise, as well as accurately describing what the paper is about. 50-100 words is generally a good guideline for length, and it should generally be only a single paragraph.

For instance, for the midterm project paper, you should mention that you used decision trees for classification, and that you used and compared both information-theoretic methods and evolutionary methods for building decision trees. You should also mention how the two techniques performed in comparison to each other. Overall, your abstract should read as a very short (but still grammatically correct) summary of what the rest of the paper is about. It should not contain details about methods or results; you'll get to those later. It should contain a brief summary of your main conclusions and/or contributions to the field. Your abstract should not be any longer than this one; as abstracts go, this is already too long.

Some publication venues expect you to have “keywords” at the end of your abstract, but you don't need to do that for this course.

NOTE: For this assignment, you are required to turn in a “partial draft” of your paper well before the final version is due. This partial draft needs to contain your *Abstract*, *Results*, and *Discussion* sections. The rest of the sections are required only in the final version. We don't want all the sections in the draft, because you wouldn't have time to

write them, and we wouldn't have time to read them and get you feedback quickly enough.

1 Introduction

In the first section of your paper, which should be entitled "Introduction," you should introduce the subject of your paper to the reader. There are several components to this. You probably want to begin by talking about what problem your work is trying to solve, and why it is a problem worth trying to solve. Explain to your reader why they should care. This is also a good time to introduce basic nomenclature related to your subject. You can assume that your reader is a reasonably well educated individual, and has some knowledge of computer science, but don't assume that they're an expert in your field. You've spent a bunch of time studying your topic (in this case, decision trees), but your reader may not have. Take the time to bring him or her up to speed; this is what the introduction is all about.

For example, you might say something like:

"The problem of classification is a fundamental one in the field of machine learning. Successful solutions to this problem form the basis for solving a wide variety of real world problems, from medical diagnosis to airport security to mobile robotics. There are a variety of techniques for attempting to solve this problem, and none of them is perfect for all situations; the famous "No free lunch" theorem[?] guarantees this. It is therefore important to compare the results of different techniques to determine what the strengths and weaknesses of various algorithms are.

Decision trees are one technique for performing classification, based on the idea of making a series of decisions about a given data vector based on simple tests applied to elements of that vector. The tree structure comes from the fact that each decision can have two outcomes, yielding a binary tree representing all possible paths through the set of decisions. Leaf nodes of the tree represent the fact that the results from the tests along the way to that leaf add up to enough information to correctly classify that example."

Of course, you wouldn't write *exactly* that, since that would be plagiarism. But it should give you the general idea. Also note that this, like the other examples in this paper, is just a short example of the kind of thing you should write. Do not use them as a guide for length, or for the minimal content that should be covered, as more is expected from you on both counts.

In terms of overall length for this paper, quality of content is more important than length; you won't be penalized for your page count in this assignment. As a rough guideline, however, your paper should probably be somewhere in the 12-18 page range (including tables, figures, and references, and using default tex margins). Shorter than that probably means that you're missing something important; longer probably means you're including stuff you don't need.

The introduction in a real paper is generally multiple pages, but yours doesn't need to be that long for this assignment. It should, however, be long enough to convey the information you need to. It should move from the initial, high level overview of the topic to a more detailed description of what other related work has been done on the topic, and any other background information needed to understand the later parts of the paper. For example, you should talk about the advantages of evolutionary algorithms in general, and what other work has been done with them. More specific details of the algorithms and techniques you used should wait for later sections.

Because your introduction discusses prior related work, you should be sure to properly cite that prior related work. If your introduction does not have at least a couple of citations, then it is either inadequate in scope, or it fails to properly cite works referenced. Don't do either. Later sections of the paper should also contain proper citations where needed, but the bulk of the citations in any given paper can usually be found in the introduction, as that is where the most related publications are discussed.

While it is important to be respectful of well-known scientists in the field, it is stylistically inappropriate in science writing to directly praise or honor them. Instead, make note of their (related) work using citations; this both indicates that you know of and respect them, and also gives someone reading your paper a place to find out more about them and their work.

Your introduction should conclude by describing the organization of the rest of the paper; for example:

“In section 2, we describe the theoretical basis for our algorithm, and in section 3, we describe the algorithm itself and our exper-

iments with it. In section 4 we present the results of our experiments, and conclude in section 5 with a discussion of what these results mean, and possible directions for future work.”

Not all papers have this sort of organizational statement, but it can be helpful due to the fact that the naming and organization of the sections in papers are not entirely standard. While almost all scientific papers begin with an introduction, then discuss the experiments and results, and finally conclude, the precise organization of that middle part can vary somewhat based on the subject matter and the preferences of the author (or sometimes the publication).

2 Learning Decision Trees

Here, you should talk in more detail about the specifics of how decision trees are implemented, as a preface to the next subsection.

2.1 Information-Theoretic Methods

Here, talk about traditional decision tree learning. You can give the generic algorithm for decision tree learning, and talk about the need for a method for choosing which variable to add to the tree next, which leads you into the next subsection. I’ve chosen not to number this subsection; you don’t need numbers everywhere, though you should number section headers. After that, it’s up to you. It’s mostly a matter of aesthetics.

Entropy and Information Gain

Here, talk about the maximum information gain heuristic. You should both give an intuitive feel for what it does and give the mathematics and theory behind it. You should also discuss why it tends to work well for decision trees, and what strengths and weaknesses it has. This should lead you into a discussion of gain ratio, which you can then define and contrast to straight information gain.

2.2 Genetic Algorithms

Here, talk about how genetic algorithms work. Give the basic algorithm, and discuss what needs to be determined (eg. encoding, fitness function, selection method, population size, etc.).

Genetic Algorithms for Decision Trees

Here, discuss the particulars of your GA. This should definitely include a discussion of your encoding, and how your mutation and crossover functions worked on it. It should also include a discussion of your fitness function, your selection method, and any other choices you made when designing your algorithm. It is very important to note that you should not only talk about what you chose, but also about how and why you chose it. This should not include how you tuned various parameters that were chosen through experimentation; that should come later, and the experiments you used should appear in your results section.

What you do need to talk about is the theoretical reasons behind your choices. For example, what makes your encoding an appropriate one for this domain? How did that impact the choice of mutation and crossover functions? How did it influence your choice of fitness function? How did you come up with your fitness function and why is it a good one for this task?

3 Algorithms and Experimental Methods

In this section, you describe any details of your algorithm that you left out of the more theoretical discussion in the previous section. You don't need to list every single parameter value, but ones that are important to your results should be discussed (eg. don't just list them, talk about why they have that value). Do not, however, discuss the implementation details of your code; the reader is probably not interested in what language you wrote in, or what data structures you used (unless your paper is about languages or datastructures, of course), and she certainly won't care what you named your variables.

You should also describe your experimental methodology; this is where you talk about your data and what you did with it. Talk about what sorts of experiments you performed, and how you validated them. For example, if you used 7-fold cross-validation, you would say that you used it, define what it is, and discuss how you implemented it. It would also be good to discuss




Figure 1: This is a caption on the figure

the strengths and weaknesses of your chosen validation method; why did you pick the one you did?

Data Sets

Here you should describe the data you used; where it came from, what it represents, what properties it has (eg. binary class? multi class? multi variate? continuous? dimensionality? number of examples? etc.). Talk about all the data sets you used. Be sure to mention and properly cite their source.

Also mention how you pre-processed the data, if you did.

4 Results

The results section should contain your results. It should *not* contain your interpretation of those results. That comes later. This section should be made up primarily of graphs and tables that show your data. You should also have a small amount of text describing what each of the tables and graphs shows, since the caption on the figures should be short. Having text describing the specifics of the experiment that lead to that particular table

	col1	col2	col3
row1	a	b	c
row2	d	e	f
row3	g		h
	i	j	k

Table 1: This is a caption on the table. Try to keep your captions short; don't put multiple paragraphs of text in here. Put the long version in the Results section, and reference the table from there.

would also be good. For example,

“Table 1 shows the average results of the three algorithms on all the data sets. The parameters used were $N = 7$, and $k = 3.27$; these parameters were found by hand, and little effort was made to tune them optimally. Each algorithm was run three times on each of the seven data sets, and the resulting accuracy scores were averaged.”

I'm not going to tell you exactly what tables or graphs you should have here, since it will depend a bit on your results. You should be sure that your results section contains sufficient data to support your conclusions about the relative strengths and weaknesses of the different algorithms. You should also be sure that your data is complete; that is, don't leave data out simply because it doesn't support the point you're trying to make.

You should also be sure that your results are clear and interpretable. Seven pages of raw binary data will do nothing to edify your reader. Similarly, a 1 inch square graph with 12 lines plotted on it will be difficult to extract meaning from, as will a graph with poor (or no) labels on the axes. Your results should be legible both on screen and in hard copy.

You don't want to present results that are just raw data, since that is hard to interpret. But you don't want to over abstract, either, since that leads to results that have little or no meaning (eg. “the average over all different data sets, algorithms, and parameters” is a completely useless statistic for comparing algorithms).

If you are trying to draw comparisons between certain things, try to make sure your results are presented in a way that allows the reader to easily

compare them. This might mean having two lines in the same graph, or it might mean having the columns you want to compare be adjacent in a table. You want it to be easy for the reader to follow whatever reasoning you make in your Conclusions section by looking at your Results. You don't want to redundancy (don't put the same data in multiple different tables), but you want the reader to be able to evaluate your conclusions without having to compare things that are three pages apart in your Results.

You should probably have a page or two of results; one or two tables are unlikely to be sufficient to describe your experiments. If they are, you need to do more experiments. On the other hand, if you have more than three pages, you're probably not doing a good job of presenting your data in a clear and concise form.

5 Discussion

The discussion section is where you discuss your interpretation of the data you presented in the results section. This is where you tell the reader how great your algorithm is, and how interesting it is that *this* performed better than *that* on some given data set. You can also speculate about causes for interesting behaviors; for example, if you think you might know why it fails so badly on some particular case, or if you have an insight into why it did well on another case. You don't want to be making wild guesses, but as long as you make it clear that you are not making claims of factual proof, you can go out on a limb a little. For example,

“In most cases, algorithm A outperforms algorithm B with a significance of 99.8%. However, as can be seen from Figure 1, when applied to the “E. E. Smith” data set, algorithm A does no better than random chance. It seems likely that the failure of algorithm A to learn is due to the extremely sparse distribution of that data set. Because of algorithm A's heavy reliance on data being densely sampled from the true underlying distribution, any sparse data set is likely to show this behavior.”

Think about the kinds of questions that were posed in the written portions of homework assignments 1 and 2; these are the kinds of things you want to think about for your Discussion.

6 Conclusions

The conclusion section should be relatively short, and should not be a summary of your paper. It should, however, bring up what you learned and what impact your results have on the rest of the field (and society as a whole, if applicable, but don't overstate the impact of what you're doing). You should conclude, and bring your paper to an end with any parting thoughts that are appropriate.

Certain types of papers can be ended with a "Summary" section instead of a "Conclusions" section, in which case you would, in fact, summarize the main points of your paper. For this paper, you should write a Conclusions section, not a Summary.

Conclusion also often contain information about what else you would like to do. Sometimes this is a separate subsection, or even a section, entitled "Future Work." The basic idea here is to talk about what the next steps to take would be. This is of benefit to others who are interested in your work and may want to help advance it. It is also a chance for you to acknowledge shortcomings in your work; since we never have infinite time to prepare a paper, there are always more experiments that would have been nice to include. If you list them as future work, then it at least makes it clear that you didn't do those things because you didn't have time, rather than because you didn't realise that they were important to do.

In your paper, you should include a brief discussion of avenues for possible future work in your Conclusions section. It should be tied in with the rest of your conclusion, and should not be an unrelated section tacked on the end (or the middle).