

# Decision Trees for Classification: An Evolutionary Approach

Ian Wilkes

April 3, 2015

## **Abstract**

Classification is one of many extremely important topics within the field of artificial intelligence, and is used in many fields from advertising to quality control. As such it is imperative that we be able to provide efficient and accurate classifiers. One such method for classification is through the use of decision trees. This paper will discuss two implementations of decision trees, one using information gain, and another which evolves to learn to classify sets of data. Their comparative performance will be discussed, as well as ways to optimize each of their individual performance.

# 1 Introduction

## 2 Learning Decision Trees

### 2.1 Information-Theoretic Methods

Entropy and Information Gain

### 2.2 Genetic Algorithms

Genetic Algorithms for Decision Trees

## 3 Algorithms and Experimental Methods

Data Sets

## 4 Results

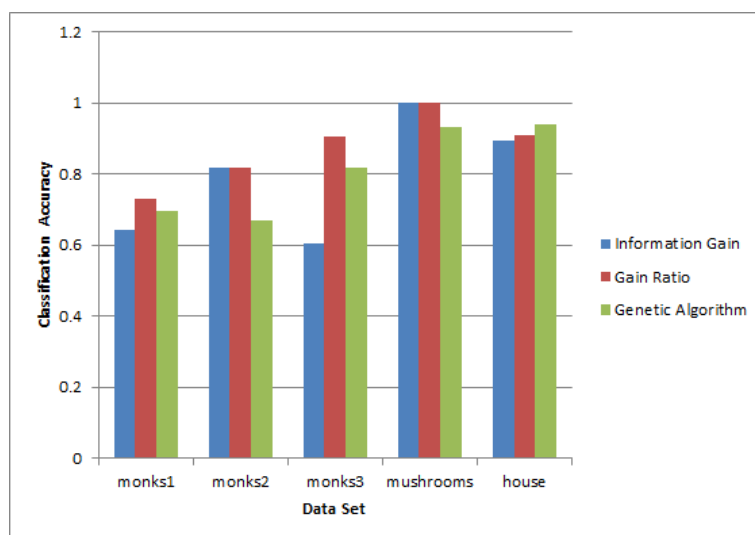


Figure 1: This figure shows the classification accuracy of the information gain based, gain ratio, and genetic decision trees. It can be seen that the gain ratio based tree seems to perform the best in almost all cases.

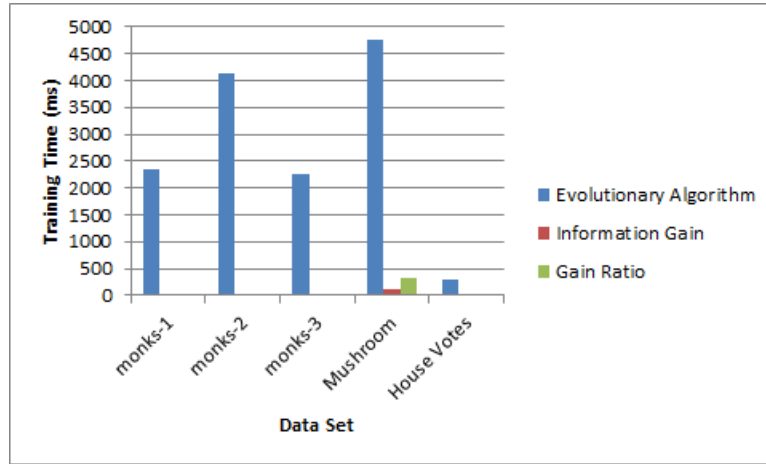


Figure 2: This figure shows the required training times of the information gain based, gain ratio, and genetic decision trees. It can be seen that the genetic tree takes much more time than either of the other trees.

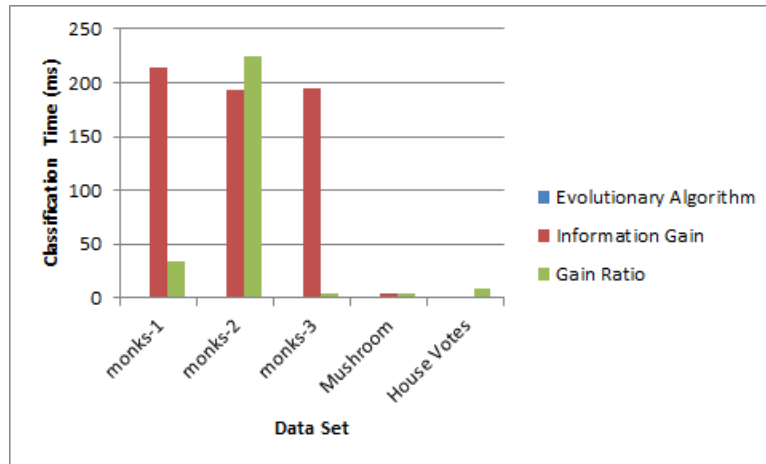


Figure 3: This figure shows the testing time of the information gain based, gain ratio, and genetic decision trees. Here, the evolutionary or genetic tree takes orders of magnitude less time than the information based trees.

For all the data shown in this section, the baseline genetic algorithm used the following parameters which were determined through minimal testing by

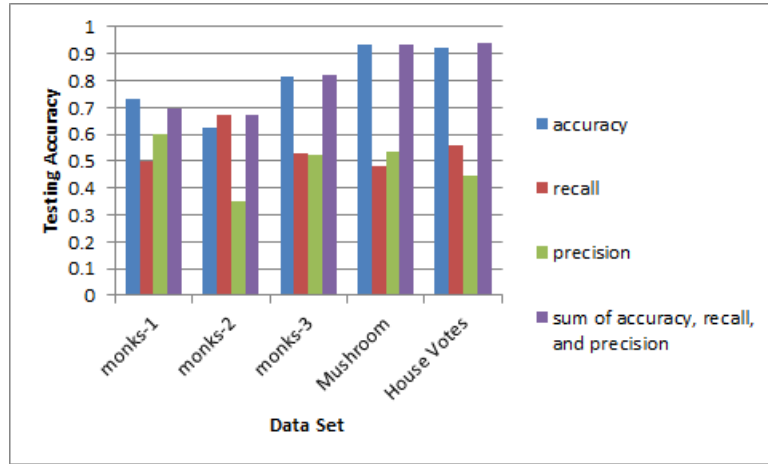


Figure 4: This figure shows the results of using different fitness functions on the classification accuracy of the genetic tree.

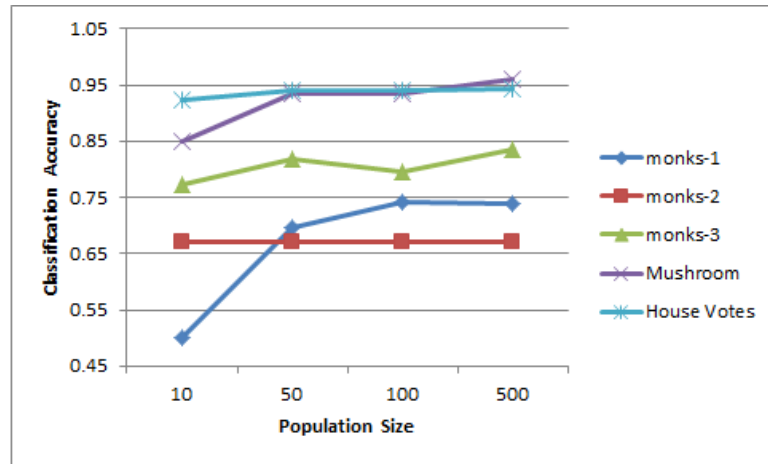


Figure 5: Here the performance of the genetic algorithm on different data sets is shown for different population sizes. For four of the five data sets, the population must be above 10 individuals to get the best classification accuracy.

hand. All parameters remain constant, except in the case where they are being examined one by one. The mutation and crossover probabilities were

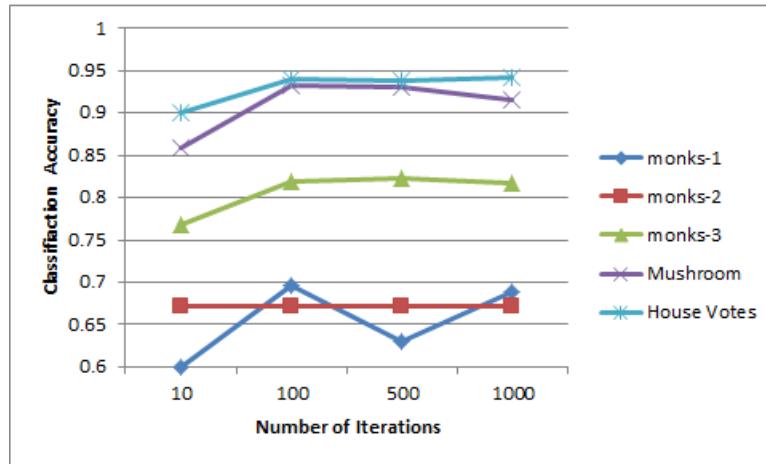


Figure 6: Here the performance of the genetic algorithm on different data sets is shown for different numbers of training iterations. For four of the five data sets, at least 100 iterations are required to get optimal classification accuracies.

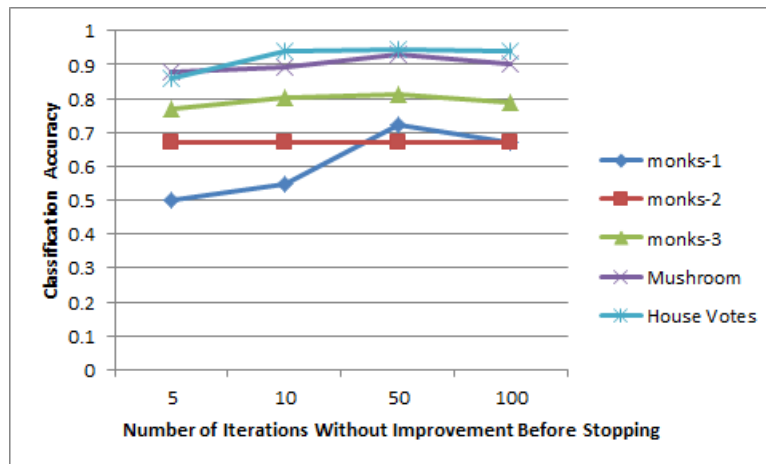


Figure 7: This plot shows the effects of the cutoff value on the classification accuracy of the genetic algorithm for different data sets. The accuracy seems to plateau after a certain cutoff value, indicating little effect.

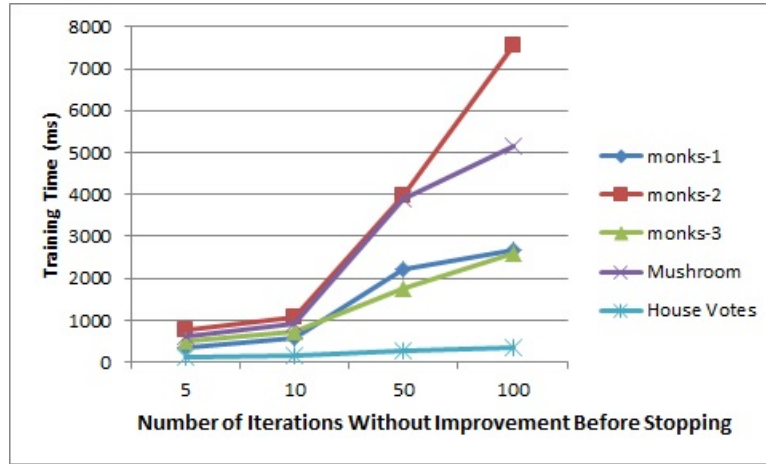


Figure 8: It can clearly be seen that the time required for training to finish increases drastically as the cutoff parameter increases for all data sets.

0.01 and 0.7 respectively. The population size was 50, and 100 iterations were performed. The default cutoff value was 50. The standard fitness function used was the sum of the classification accuracy, recall and precision for each iteration. Finally, all data presented here is the average of five trials for each data point.

From figure 1 we can see that the performances of the genetic, information gain and gain ratio based trees are all somewhat comparable. The gain ratio tree seems to perform the best of the three in all but the house votes data set. It drastically out performs the information gain based tree on the monks-3

	col1	col2	col3
row1	a	b	c
row2	d	e	f
row3	g		h
	i	j	k

Table 1: This is a caption on the table. Try to keep your captions short; don't put multiple paragraphs of text in here. Put the long version in the Results section, and reference the table from there.

data set, which is also beaten by the genetic algorithm.

Moving on to figure 2, it is plain to see that the evolutionary algorithm takes much longer to train than the traditional decision trees, neither of which ever takes even 500ms to train, while the genetic tree takes almost 5000ms on the mushroom data set.

The opposite trend can be seen in the classification times of the testing sets, which can be observed in figure 3. The information gain and gain ratio trees take orders of magnitude more time to classify their data on the monks-1, monks-2, and monks-3 data sets than the genetic tree.

Figures 4 through 10 show the results from various experiments performed to try to optimize the genetic tree shown above. The first of these is figure 4 which shows the effect on classification accuracy of various fitness functions. It can clearly be seen that by not using the accuracy as one of the components of the fitness function, the resulting classification accuracy is greatly reduced, as seen in the cases where recall and precision were used as the fitness metrics.

Another parameter which was varied was the population size within the genetic algorithm based tree. From figure 5, it can be seen that below a certain population threshold the classification performance decreases, yet above that threshold, an increase in population size does not necessarily result in improved classification performance. The same trend can be seen for the number of iterations used to train the tree in figure 6. Other than the monks-2 data set, all data sets showed a classification accuracy increase by increasing the number of iterations from 10 to 100.

The effects of mutation and crossover probabilities can be seen in figures 7 and 8 respectively. In both cases, the change in the probabilities seem to cause different results for the different data sets. Increasing the mutation seems to hurt the Mushroom accuracy, while it helps the monks-1 accuracy going from 0.03 to 0.05. Similarly for crossover, accuracy is highest for monks-3 when the crossover probability is 0.9, but this value causes the accuracy for both the monks-1 and mushroom data sets to decrease significantly.

Finally, the influence of the cutoff parameter, which stops iteration if no improvements have been found after a given number of iterations, is shown in figures 9 and 10. In figure 9, the effect on accuracy can be seen to be minimal for high values of the cutoff parameter, as the performance for most data sets stays fairly constant. Figure 10 shows that while the performance may not change, the computation time to train the GA increases substantially.

## 5 Discussion

The discussion section is where you discuss your interpretation of the data you presented in the results section. This is where you tell the reader how great your algorithm is, and how interesting it is that *this* performed better than *that* on some given data set. You can also speculate about causes for interesting behaviors; for example, if you think you might know why it fails so badly on some particular case, or if you have an insight into why it did well on another case. You don't want to be making wild guesses, but as long as you make it clear that you are not making claims of factual proof, you can go out on a limb a little. For example,

“In most cases, algorithm A outperforms algorithm B with a significance of 99.8%. However, as can be seen from Figure ??, when applied to the “E. E. Smith” data set, algorithm A does no better than random chance. It seems likely that the failure of algorithm A to learn is due to the extremely sparse distribution of that data set. Because of algorithm A's heavy reliance on data being densely sampled from the true underlying distribution, any sparse data set is likely to show this behavior.”

Think about the kinds of questions that were posed in the written portions of homework assignments 1 and 2; these are the kinds of things you want to think about for your Discussion.

## 6 Conclusions

The conclusion section should be relatively short, and should not be a summary of your paper. It should, however, bring up what you learned and what impact your results have on the rest of the field (and society as a whole, if applicable, but don't overstate the impact of what you're doing). You should conclude, and bring your paper to an end with any parting thoughts that are appropriate.

Certain types of papers can be ended with a “Summary” section instead of a “Conclusions” section, in which case you would, in fact, summarize the main points of your paper. For this paper, you should write a Conclusions section, not a Summary.



Conclusion also often contain information about what else you would like to do. Sometimes this is a separate subsection, or even a section, entitled “Future Work.” The basic idea here is to talk about what the next steps to take would be. This is of benefit to others who are interested in your work and may want to help advance it. It is also a chance for you to acknowledge shortcomings in your work; since we never have infinite time to prepare a paper, there are always more experiments that would have been nice to include. If you list them as future work, then it at least makes it clear that you didn’t do those things because you didn’t have time, rather than because you didn’t realise that they were important to do.

In your paper, you should include a brief discussion of avenues for possible future work in your Conclusions section. It should be tied in with the rest of your conclusion, and should not be an unrelated section tacked on the end (or the middle).

## References